The Inevitable Limits of Adaptive Systems: a Structural Theory of Drift, Entropy, and Survival Horizons

Maksim Barziankou, Poznan 2026

Abstract

Adaptive systems are commonly analyzed through the lenses of optimization, control accuracy, stability, or reward maximization. Such approaches implicitly assume that sufficiently precise regulation can, in principle, preserve system viability indefinitely. This work rejects that assumption.

We present a theory of inevitable limits governing the long-horizon existence of adaptive systems. We show that drift, entropy growth, irreversibility, loss of causal trace, structural asymmetry, latent residual accumulation, and internal time depletion are not implementation flaws but unavoidable consequences of directed interaction under finite resources. These phenomena form a tightly coupled chain of inevitabilities that constrains survival independently of action-level correctness or performance quality.

The result is a fundamental decoupling between optimization success and identity preservation. Long-horizon survival is shown to be governed by structural limits orthogonal to control accuracy. Any architecture that ignores these limits is provably incomplete with respect to viability and identity continuity.

(Expansion) The central contribution of this work is not the identification of any single failure mode, but the demonstration that failure is structurally delayed. Systems do not collapse when they err; they collapse when accumulated structural cost exhausts the remaining capacity for identity-preserving evolution. This distinction is critical: it shifts the analysis of adaptive systems from what they do to what they irreversibly consume by doing it.

Introduction

Most theories of adaptive systems are implicitly optimistic. They assume that failure arises from insufficient sensing, imperfect modeling, limited computation, or suboptimal control. The dominant intuition is corrective: if drift appears, compensate it; if instability arises, regulate it; if performance degrades, optimize harder.

This work advances a different claim:

Some forms of degradation are not correctable, not because of technical limitations, but because they are structurally inevitable.

The goal of this paper is not to propose a better controller, learning rule, or safety mechanism. Instead, we characterize a class of limits that apply to all adaptive systems operating under directed interaction, regardless of substrate, intelligence, or design.

(Expansion) In particular, we challenge the hidden premise that perfect regulation implies indefinite survival. Even under idealized assumptions—unbounded precision, arbitrarily fast correction, and flawless local decision-making—long-horizon identity cannot be preserved. The reason is not error, but history. Every adaptive act embeds irreversible structure into the system's future, even when that act is locally optimal.

Directed Interaction and the Emergence of Drift

Consider any adaptive system interacting with an environment in a directed manner—that is, the system exerts influence while receiving feedback under asymmetric coupling.

Under such interaction, perfect symmetry is impossible. Each action leaves a residual mismatch between intended structure and realized consequence. This mismatch cannot be fully eliminated locally without additional intervention.

We define drift as the accumulation of such irreducible residual asymmetries over time.

Claim 1 (Inevitability of Drift). Under sustained directed interaction with a changing environment, accumulated drift is unavoidable.

Drift is not noise, error, or instability. It is the structural trace of history.

(Expansion) Crucially, drift persists even when actions are admissible, safe, and correct. It arises not from mistakes, but from directionality itself. Any system that imposes preferred directions of change—explicitly or implicitly—introduces antisymmetric deformation between intention and outcome. This deformation cannot be symmetrically unwound without cost, because the environment does not replay the past.

Entropy as Loss of Reversibility

Drift alone does not yet define a limit. The limit arises when we consider reversibility.

Compensating drift requires intervention. Intervention itself leaves a trace. The path of compensation is not the inverse of the path of deviation. As a result, even perfectly restoring admissible behavior does not restore the prior structural state.

This leads to entropy growth, understood here not thermodynamically, but structurally: as loss of the ability to uniquely reconstruct prior states from present structure.

Claim 2 (Inevitability of Entropy Growth). If drift is non-zero, entropy growth is unavoidable. Structural reversibility cannot be preserved indefinitely under compensation.

Thus, every act of correction trades local stability for global irreversibility.

(Expansion) This notion of entropy does not depend on heat, randomness, or statistical disorder. It refers instead to causal compression: multiple distinct histories collapse into indistinguishable present states. Once this occurs, no amount of local regulation can recover what has been structurally forgotten.

Irreversibility and Regime Boundaries

Entropy implies irreversibility. Over time, systems cross boundaries beyond which return to a prior organizational regime is impossible—even if behavior is restored.

Such boundaries define irreversible regime transitions.

Claim 3 (Inevitability of Irreversible Transitions). There exist regime transitions after which restoration of admissible behavior does not imply restoration of prior identity.

This establishes a fundamental distinction between working again and being the same system again.

(Expansion) This distinction is routinely ignored in engineering practice, where recovery is equated with success. Yet many systems continue operating in a post-identity regime—functionally acceptable, but structurally altered. The theory presented here treats such transitions not as anomalies, but as expected milestones in long-horizon operation.

Loss of Causal Trace

As drift accumulates and entropy grows, the system progressively loses information about how it arrived at its current state. Causes collapse into effects; history compresses.

The system may still function, but it no longer knows why it functions as it does.

Claim 4 (Inevitability of Causal Trace Loss). Long-horizon operation necessarily erodes internal access to causal history, even when external performance remains stable.

This loss prevents precise diagnosis of degradation from within the causal decision process.

(Expansion) At this stage, the system becomes operationally opaque to itself. Decisions are made without reliable access to the structural reasons they remain admissible. This is not a monitoring failure, but a consequence of accumulated irreversibility.

Structural Asymmetry and Latent Residuals

Compensation is not neutral. Each corrective action redistributes burden unevenly across structural degrees of freedom. Over time, this produces structural asymmetry.

Some mismatches remain latent: they do not appear as immediate errors or constraint violations, yet they accumulate as unresolved structural debt.

Claim 5 (Inevitability of Latent Residual Accumulation). Even under bounded error and admissible trajectories, latent structural residuals accumulate under sustained interaction.

This explains why systems often fail "suddenly" after long periods of apparent stability.

(Expansion) Latent residuals are invisible to performance metrics by construction. They do not announce themselves until structural capacity is exhausted. When failure finally manifests, it appears abrupt only because its causes were distributed across time.

Internal Time and Survival Horizon

We define internal time as the remaining structural viability budget of a system—the amount of deformation it can absorb before identity collapses.

Internal time is not synchronized with external clocks, nor is it reducible to performance metrics.

Claim 6 (Inevitability of Internal Time Depletion). Internal time decreases monotonically under drift and latent residual accumulation, regardless of instantaneous correctness.

Once internal time is exhausted, continued operation necessarily violates identity constraints.

(Expansion) Internal time formalizes the intuition that systems "age" structurally. Two systems may be equally performant at the same external time, yet possess radically different remaining viability horizons due to differences in accumulated structural cost.

Performance–Identity Decoupling

The preceding claims lead to a central result:

Claim 7 (Inevitability of Performance–Identity Decoupling). There exists an interval during which an adaptive system preserves admissible behavior and stable performance while its identity has already been structurally compromised.

This interval is not anomalous; it is unavoidable.

Failure, therefore, is not the moment performance collapses, but the delayed consequence of earlier structural exhaustion.

(Expansion) This decoupling invalidates guarantees based solely on reward, error, or stability. Systems can pass every test that matters—until they suddenly cannot. The tests were never designed to observe identity depletion.

Non-Causal Observability as a Necessity

Because the drivers of long-horizon degradation are cumulative, delayed, and non-local, embedding their observation into the causal decision loop creates distortion.

Either observation alters behavior, or degradation remains invisible.

Claim 8 (Inevitability of Non-Causal Observability). Any system that seeks to detect long-horizon degradation must employ a non-causal observational dimension, decoupled from action selection.

This is not an architectural choice but a structural requirement.

(Expansion) Observation must occur about the system, not through it. Any attempt to fold such observation into the action loop reintroduces the very coupling that accelerates depletion.

The Chain of Inevitabilities

The results can be summarized as a single chain:

Directed interaction → Drift → Entropy growth → Irreversibility → Loss of causal trace → Structural asymmetry → Latent residual accumulation → Internal time depletion → Identity–performance decoupling → Coherence collapse before failure

The chain is structural, not empirical; it holds independently of implementation or intelligence level. Breaking any link requires violating a prior assumption (e.g., infinite resources, perfectly reversible interaction, or absence of directionality).

(Expansion) The chain is not speculative. Each link constrains the next. Removing any element collapses the descriptive power of the theory and reintroduces implicit optimism.

Implications

The results presented here impose principled limits on a broad class of approaches to adaptive system design.

In particular, action-centric safety frameworks, reward-maximization guarantees, perpetual learning paradigms, and claims of indefinite alignment through optimization are structurally incomplete with respect to long-horizon viability and identity continuity.

No amount of local correctness, bounded error, or stable performance can compensate for accumulated structural exhaustion. Survival is therefore not an optimization objective, but a finite budget governed by irreversible structural consumption.

(Expansion) Design, therefore, must shift from preventing failure to managing inevitability. The relevant question is not "How do we avoid drift?" but "How do we allocate the remaining internal time responsibly?"

Conclusion

Adaptive systems do not fail because they are insufficiently intelligent. They fail because they exist in time.

Drift cannot be eliminated without cost. Entropy cannot be reversed through control. Identity cannot be preserved indefinitely under directed interaction, regardless of local correctness or optimization quality.

The central challenge of long-horizon adaptive systems is therefore not the prevention of degradation, but the recognition, regulation, and principled management of inevitable limits. Survival is not a function of intelligence alone, but of how structural consumption is acknowledged and bounded.

Any theory of adaptive intelligence that does not explicitly account for drift, entropy, irreversibility, and internal time is structurally incomplete, regardless of its empirical performance.

Meta-conclusion

These limits are not artifacts of insufficient design, intelligence, or optimization capacity, but arise from the mere fact of persistent existence under directed interaction.

MxBv, 2025 - 2026

DOI: 10.5281/zenodo.18208692